

# Model Zoo 3


The textbook has detailed explanations of all of the models we've covered in the model zoo in the online supplements located at

<https://amlbook.com/eChapters.html>

The online supplements are all password protected with password **Paraskavedekatriaphobia**. For the most part, these online supplements are technical and provide relatively little insight into practical applications. The section on nearest neighbor algorithms (e-Chapter 6), however, is relatively readable.

## Nearest Neighbor Methods

**Problem 1.** Describe the  $k$ -nearest neighbor classifier.



$k \uparrow \Rightarrow$  noise is less of a problem

## Runtime

Brute force: runtime  $\Theta(dN)$

scikit learn: ball tree, kdtree  $O(2^d N)$

clever data structure:

goal: each search take time  $O(\log N)$

practice:  $O(2^d \log N)$

# curse of dimensionality

**Fact 1.** The in sample error of 1-nearest neighbor is always 0.

$$\Rightarrow d_{VC} = \infty$$

$$\Rightarrow |E_{in} - E_{out}| = O\left(\sqrt{\frac{d_{VC}}{N}}\right)$$

is useless for nearest neighbor

**Theorem 1** (informal). Let  $h$  be the 1-NN hypothesis. Then for "well behaved" data distributions, we have with high probability that

$$E_{out}(h) \leq 2E_{out}(f) + 4\sqrt{d}N^{-\frac{1}{d+1}} \quad \text{exp. in } d \quad (1)$$

If  $h$  is the  $k$ -nearest neighbor classifier, then with high probability,

$$E_{out}(h) \geq (1 + 1/k)E_{out}(f). \quad (2)$$

*Proof.* See Theorem 19.3 of *Understanding Machine Learning: From Theory to Algorithms*.  $\square$

$$\textcircled{1} E_{out}(h) \approx 2E_{out}(f)$$

if  $E_{out}$  is very small  $\Rightarrow$  1NN might be good  
if " large  $\Rightarrow$  1NN definitely bad

$$\textcircled{2} \text{gen}_{err} = \sqrt{\frac{d_{VC}}{N}} \sim \text{linear} = d$$

$$N^{-\frac{1}{d+1}} = \epsilon$$

$$N = \left(\frac{1}{\epsilon}\right)^{d+1}$$

exponential in  $d$

As  $d \uparrow \Rightarrow N \uparrow$  exponentially

practice  
rule of thumb

$d \leq 10 \Rightarrow$  okay

$d > 10 \Rightarrow$  very concerned


curse of dimensionality

**Problem 2.** Recall that the ImageNet dataset has an estimated Bayes Error around 5%. (This estimate follows from Karpathy's observed human error rate of 5% and the observed labelled errors from the automatic collection process.) What is the best out of sample error we can expect from a 1-NN classifier?

$\approx 10\%$

**Problem 3.** Recall that the MNIST digit classification dataset has  $N = 60000$  training image, and each image has  $d = 28 \times 28 = 784$  dimensions. There are 13 known mislabeled images. (You can see them at <https://labelerrors.com/>.)

1. Why 1-NN is a bad choice of algorithm for this dataset?
2. Would it make more sense to use PCA or the polynomial feature map on this dataset before performing 1-NN?




## Support Vector Machines (SVMs)

There are three standard and equivalent interpretations of SVMs:

1. The hypothesis class  $\mathcal{H}_{\text{perceptron}}$  with the hinge loss  $\ell(\mathbf{x}; \mathbf{w}) = \max(1 - \mathbf{w}^T \mathbf{x}, 0)$ .
2. The large margin classifier.
3. A smooth generalization of the k-nearest neighbor algorithm.

Convex, surrogate, most robust



all valid perceptrons  
PLA might return any  
of these lines

**Note 1.** The following theorem is copied directly from Theorem 8.5 in Chapter 8 of the *Learning from Data* textbook. Most of this chapter is more technical than needed for this course, but the explanation of the theorem is straightforward and worth reading.

**Theorem 2.** Suppose the input space is the ball of radius  $R$  in  $\mathbb{R}^d$ . That is,  $\|\mathbf{x}\|_2 \leq R$ . Then,

$$d_{\text{VC}} \leq \boxed{\lceil R^2/\rho^2 \rceil + 1.} \quad \text{no dep. on } d \quad (3)$$

SVM is a linear classifier, it also

satisfies  $d_{\text{VC}} \leq d+1$

**Problem 4.** Describe the *dual learning problem* and the *kernel trick*.

Common sample kernel functions include:

kernel name	$K(\mathbf{x}_1, \mathbf{x}_2)$	feature dimensions ( $\tilde{d}$ )
linear	$\mathbf{x}_1^T \mathbf{x}_2$	$d$
polynomial	$(\gamma \mathbf{x}_1^T \mathbf{x}_2 + r)^Q$	$\Theta(d^Q)$
gaussian	$\exp(-\gamma \ \mathbf{x}_1 - \mathbf{x}_2\ _2^2)$	$\infty$
sigmoid	$\tanh(\gamma \mathbf{x}_1^T \mathbf{x}_2 + r)$	$\infty$